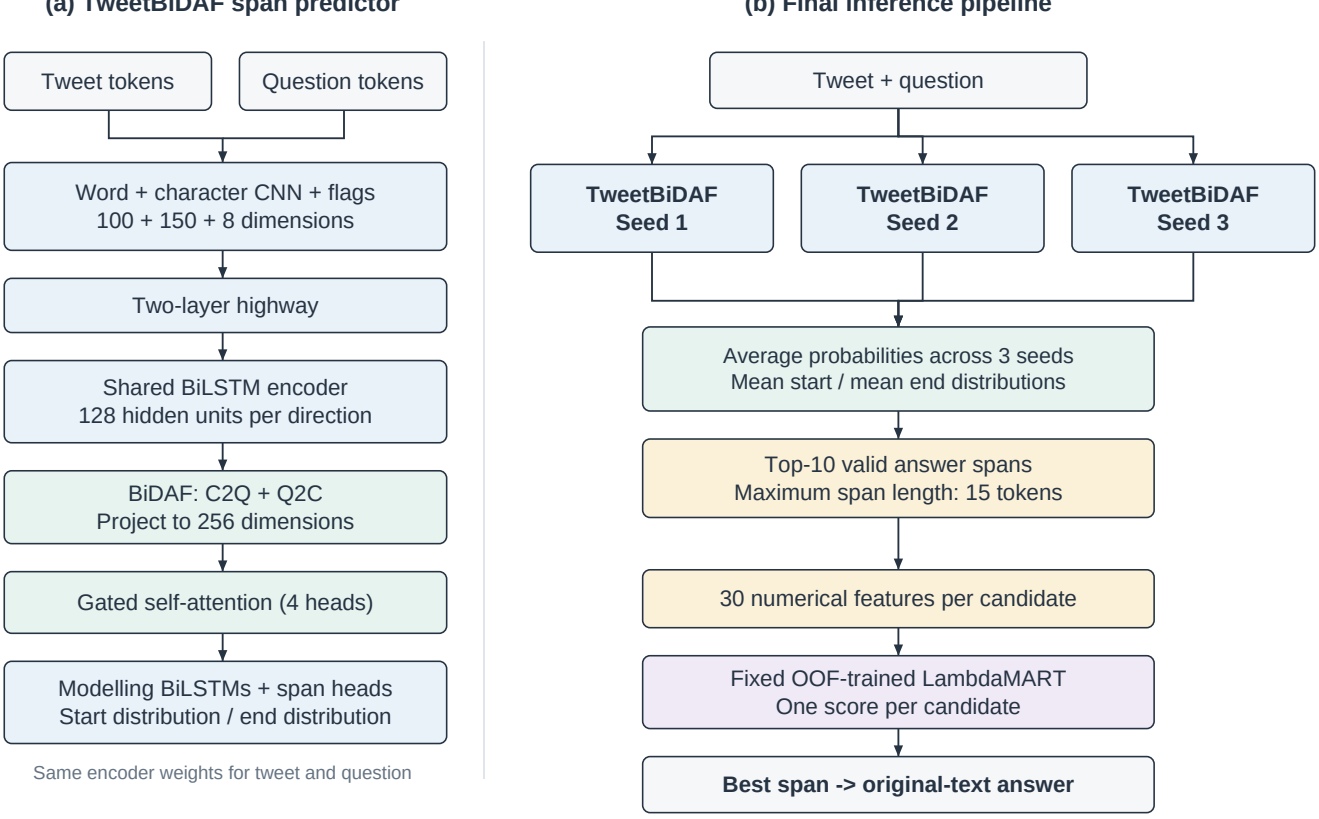
(a) TweetBiDAF span predictor
Tweet tokens
Question tokens
Word + character CNN + flags
100 + 150 + 8 dimensions
Two-layer highway
Shared BiLSTM encoder
128 hidden units per direction
BiDAF: C2Q + Q2C
Project to 256 dimensions
Gated self-attention (4 heads)
Modelling BiLSTMs + span heads
Start distribution / end distribution
Same encoder weights for tweet and question
(b) Final inference pipeline
Tweet + question
TweetBiDAF Seed 1
TweetBiDAF Seed 2
TweetBiDAF Seed 3
Average probabilities across 3 seeds
Mean start / mean end distributions
Top-10 valid answer spans
Maximum span length: 15 tokens
30 numerical features per candidate
Fixed OOF-trained LambdaMART
One score per candidate
Best span -> original-text answer