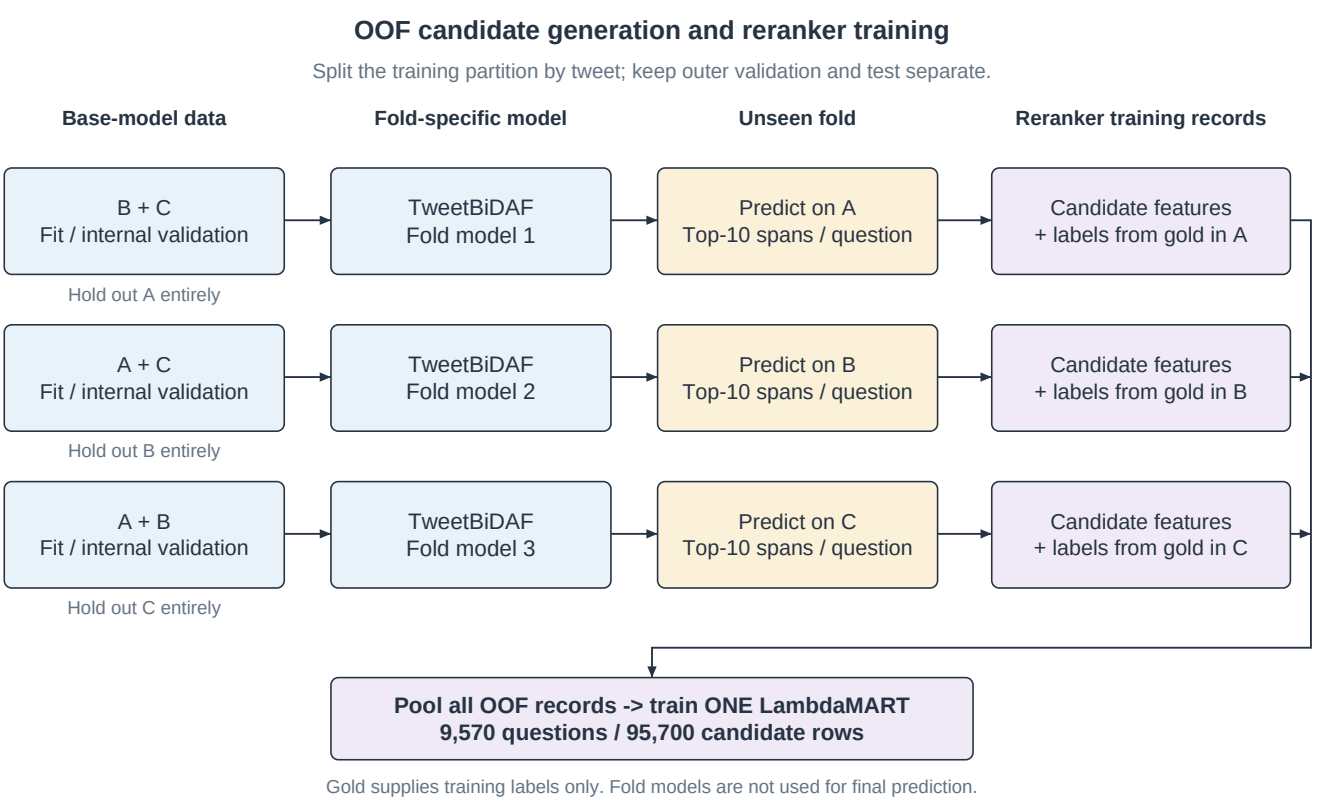
OOF candidate generation and reranker training
Split the training partition by tweet; keep outer validation and test separate.
Base-model data
Fold-specific model
Unseen fold
Reranker training records
B + C
Fit / internal validation
TweetBiDAF
Fold model 1
Predict on A
Top-10 spans / question
Candidate features
+ labels from gold in A
Hold out A entirely
A + C
Fit / internal validation
TweetBiDAF
Fold model 2
Predict on B
Top-10 spans / question
Candidate features
+ labels from gold in B
Hold out B entirely
A + B
Fit / internal validation
TweetBiDAF
Fold model 3
Predict on C
Top-10 spans / question
Candidate features
+ labels from gold in C
Hold out C entirely
Pool all OOF records -> train ONE LambdaMART
9,570 questions / 95,700 candidate rows
Gold supplies training labels only. Fold models are not used for final prediction.